

```
import numpy as np
import matplotlib.pyplot as plt
from matplotlib import rcParams
import os
```

```
# pdf of exponential distribution
```

```
def exp_dist_PDF(x, l):
    # lambda e^(-lambda x)
    return l*np.exp(-l*x)
```

```
def plot_mixture_pdf(pi_1, pi_2, l_1, l_2, name, ylab, iters=False):
```

```
    # generate samples
```

```
    num_samples = 10000
    exp_mixture = np.zeros(shape=(num_samples, 2))
```

```
    for i in range(num_samples):
```

```
        # rescale i
```

```
        x_input = i / 1000
```

```
        # mixture PDF
```

```
        mix_y = pi_1 * exp_dist_PDF(x_input, l_1) + pi_2 * exp_dist_PDF(x_input, l_2)
        exp_mixture[i] = [x_input, mix_y]
```

```
    # save PDF data
```

```
    plt.plot(exp_mixture[:, 0], exp_mixture[:, 1])
```

```
    plt.title(name)
```

```
    plt.xlabel(r'$x$')
```

```
    plt.ylabel(ylab)
```

```
    if iters:
```

```
        plt.savefig("plots/iters/" + name + ".png")
```

```
    else:
```

```
        plt.savefig("plots/" + name + ".png")
```

```
    plt.clf()
```

```
    plt.close()
```

```
def plot_log_likelihood_history(LL_hist):
```

```
    x_axis = np.arange(LL_hist.shape[0])
```

```
    plt.plot(x_axis, LL_hist)
```

```
    plt.xlabel("Iterations")
```

```
    plt.ylabel(r'$\log f(X ; \hat{\theta})$')
```

```
    plt.title("Log Likelihood")
```

```
    plt.savefig("plots/LogLikelihood.png")
```

```
    plt.clf()
```

```
    plt.close()
```

```
def E_step(data, pi_1, pi_2, l1, l2):
```

```
    # compute conditional expectation for each mixture
```

```
    Q_theta = np.zeros(shape=(data.shape[0], 2))
```

```
    for i in range(data.shape[0]):
```

```
        # get data point
```

```
        xi = data[i]
```

```
        # compute each expectation
```

```
        qi1_num = pi_1 * exp_dist_PDF(xi, l1)
```

```
        qi2_num = pi_2 * exp_dist_PDF(xi, l2)
```

```
        qi_denom = qi1_num + qi2_num
```

```
        # add it to expectation buffer
```

```
        Q_theta[i][0] = qi1_num / qi_denom
```

```
        Q_theta[i][1] = qi2_num / qi_denom
```

```
    # return the conditional expectation
```

```
    return Q_theta
```

```

def M_step(data, Q_theta):

    # update parameters
    l1_update_num = 0
    l1_update_denom = 0

    l2_update_num = 0
    l2_update_denom = 0

    pi_1_update_num = 0
    pi_1_update_denom = 0

    pi_2_update_num = 0
    pi_2_update_denom = 0

    # for each data point
    for i in range(data.shape[0]):
        # compute parameter update for each data point
        pi_1_update_num += Q_theta[i][0]
        pi_2_update_num += Q_theta[i][1]

        pi_1_update_denom += Q_theta[i][0] + Q_theta[i][1]
        pi_2_update_denom += Q_theta[i][0] + Q_theta[i][1]

        l1_update_num += Q_theta[i][0]
        l2_update_num += Q_theta[i][1]

        l1_update_denom += Q_theta[i][0]*data[i]
        l2_update_denom += Q_theta[i][1]*data[i]

    # combine all update rules to get the parameter update
    pi_1_update = pi_1_update_num / pi_1_update_denom
    pi_2_update = pi_2_update_num / pi_2_update_denom

    l1_update = l1_update_num / l1_update_denom
    l2_update = l2_update_num / l2_update_denom
    # return all parameters
    return pi_1_update, pi_2_update, l1_update, l2_update

```

```

def check_early_termination(theta_hist, idx):
    # make sure we're not checking this at the first iteration
    # otherwise we get an out of bounds error
    if idx > 0:
        # if every element in the previous iteration is exactly the same
        # end the EM algorithm at it reached a local or global minima
        if np.array_equal(theta_hist[idx], theta_hist[idx-1]):
            return True
        # otherwise don't terminate and keep going
    else:
        return False

```

```

def current_log_likelihood(data, pi_1, pi_2, l1, l2):
    # compute the log likelihood of the current estimate
    LL = 0
    # for every datapoint
    for d in data:
        # sum the log likelihood at this datapoint
        LL += np.log(pi_1*exp_dist_PDF(d, l1) + pi_2*exp_dist_PDF(d, l2))
    return LL

```

```
def EM(data):
```

```
    # randomly initialize guesses
```

```
    l1 = np.random.rand()*5.0
```

```
    l2 = np.random.rand()*5.0
```

```
    pi_1 = np.random.rand()
```

```
    pi_2 = 1 - pi_1
```

```
    # fixed number of steps if our estimate doesn't converge
```

```
    num_steps = 100000
```

```
    # history of estimates
```

```
    theta_history = np.zeros(shape=(num_steps+1, 4))
```

```
    log_likelihood_history = []
```

```
    # save first guess in history buffer
```

```
    theta_history[0] = [l1, l2, pi_1, pi_2]
```

```
    for i in range(num_steps):
```

```
        # Compute log likelihood for current estimate
```

```
        log_likelihood_history.append(current_log_likelihood(data, pi_1, pi_2, l1, l2))
```

```
        # E step
```

```
        Q_theta = E_step(data, pi_1, pi_2, l1, l2)
```

```
        # M Step
```

```
        pi_1, pi_2, l1, l2 = M_step(data, Q_theta)
```

```
        # append data to our estimation history
```

```
        theta_history[i+1] = [l1, l2, pi_1, pi_2]
```

```
        # check for early termination
```

```
        if check_early_termination(theta_history, i):
```

```
            # early termination, slice history array and return
```

```
            theta_history = theta_history[:i+2]
```

```
            break
```

```
    return theta_history, np.asarray(log_likelihood_history)
```

```
# rejection sampling from mixture PDF
```

```
def rejection_sampling(l_1, l_2, pi_1, pi_2, num_samples = 20):
```

```
    # count number of accepted samples
```

```
    num_accepted = 0
```

```
    # max of our pdf
```

```
    max_pdf = pi_1 * exp_dist_PDF(0, l_1) + pi_2 * exp_dist_PDF(0, l_2)
```

```
    # store all samples here
```

```
    samples = []
```

```
    # keep iterating until we get the desired number of samples
```

```
    while num_accepted != num_samples:
```

```
        # uniform sample from a support
```

```
        x_i = np.random.rand()*10
```

```
        # get the pdf of that sample
```

```
        pdf_y = pi_1 * exp_dist_PDF(x_i, l_1) + pi_2 * exp_dist_PDF(x_i, l_2)
```

```
        # uniform sample from our support of PDFs
```

```
        u_i = np.random.rand()*max_pdf
```

```
        # if our random sample is less than the sampled pdf
```

```
        # accept the sample
```

```
        if u_i < pdf_y:
```

```
            samples.append(x_i)
```

```
            num_accepted += 1
```

```
    return np.asarray(samples)
```

```
if __name__ == "__main__":
```

```
    # Set to true for plot of each iteration
```

```
    DETAILED_OUTPUT = False
```

```

# create directories for saving figures
if not os.path.exists("plots"):
    os.makedirs("plots")

if DETAILED_OUTPUT:
    if not os.path.exists("plots/iters"):
        os.makedirs("plots/iters")

# fix seed for reproducibility
np.random.seed(999)

# plotting params
rcParams.update({'figure.autolayout': True})

# true parameters
true_l1 = 1
true_l2 = 3
true_pi_1 = 0.25
true_pi_2 = 1.0 - true_pi_1

# plot the mixture pdf
plot_mixture_pdf(l_1=true_l1, l_2=true_l2, pi_1=true_pi_1, pi_2=true_pi_2
                 , name="ExponentialMixture", ylab=r'$f(x; \hat{\theta})$')

# generate n observations
data = rejection_sampling(l_1=true_l1, l_2=true_l2, pi_1=true_pi_1, pi_2=true_pi_2)

# run EM
theta_history, LL_history = EM(data)

# grab the last estimate of our parameters
theta_history_last = theta_history[-1]

# grab each parameters separately to plot pdf
estimated_l1 = theta_history_last[0]
estimated_l2 = theta_history_last[1]
estimated_pi_1 = theta_history_last[2]
estimated_pi_2 = theta_history_last[3]

# plot final estimated parameters pdf
plot_mixture_pdf(l_1=estimated_l1, l_2=estimated_l2, pi_1=estimated_pi_1, pi_2=estimated_pi_2
                 , name="EstimatedExponentialMixture", ylab=r'$f(x; \hat{\theta})$')

if DETAILED_OUTPUT:
    # plot estimated parameter pdf for some iterations
    # save every x steps
    save_every = 1
    for i in range(theta_history.shape[0]):
        estimate_at_i = theta_history[i]
        if (i % save_every) == 0:
            estimated_l1 = estimate_at_i[0]
            estimated_l2 = estimate_at_i[1]
            estimated_pi_1 = estimate_at_i[2]
            estimated_pi_2 = estimate_at_i[3]
            # plot estimated parameters pdf for current step
            plot_mixture_pdf(l_1=estimated_l1, l_2=estimated_l2, pi_1=estimated_pi_1, pi_2=estimated_pi_2
                             , name="EstimatedExponentialMixture_Iteration"+str(i), ylab=r'$f(x; \hat{\theta})$',
                             iters=True)

plot_log_likelihood_history(LL_history)
np.savetxt("plots/theta_history.txt", theta_history)

```